

AlexNet using CIFAR-10

April 3, 2025

```
[3]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout, BatchNormalization
from tensorflow.keras.datasets import cifar10
from tensorflow.keras import optimizers
import numpy as np
from sklearn.metrics import classification_report, confusion_matrix,
↳precision_score, recall_score, f1_score
import matplotlib.pyplot as plt
```

```
[5]: #Loading the data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
[7]: from tensorflow.keras.utils import to_categorical

y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

```
[9]: #Normalising the data to get the vales in the range [0,1]
x_train, x_test = x_train.astype("float32") / 255.0, x_test.astype("float32") /
↳255.0
```

```
[11]: train_datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
    zoom_range=0.1,
)
```

```
[13]: test_datagen = ImageDataGenerator()
```

```
[15]: #Generating images for training and testing in the model
train_generator = train_datagen.flow(x_train, y_train, batch_size=128)
test_generator = test_datagen.flow(x_test, y_test, batch_size=128)
```

```
[17]: model = Sequential([
    Conv2D(96, (3, 3), strides=(1, 1), activation='relu', padding='same',
    ↪input_shape=(32, 32, 3)),
    BatchNormalization(),
    MaxPooling2D(pool_size=(3, 3), strides=(2, 2)),
    Conv2D(256, (3, 3), padding='same', activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(3, 3), strides=(2, 2)),
    Conv2D(384, (3, 3), padding='same', activation='relu'),
    BatchNormalization(),
    Conv2D(384, (3, 3), padding='same', activation='relu'),
    BatchNormalization(),
    Conv2D(256, (3, 3), padding='same', activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(3, 3), strides=(2, 2)),
    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.3),
    Dense(256, activation='relu'),
    Dropout(0.3),
    Dense(10, activation='softmax')
])

optimizer=tf.keras.optimizers.AdamW(learning_rate=1e-3, weight_decay=1e-4,
    ↪clipvalue=1.0)
model.compile(optimizer=optimizer,loss='categorical_crossentropy',
    metrics=['accuracy'])
```

```
/opt/anaconda3/envs/tf-metal/lib/python3.9/site-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
2025-04-03 08:55:48.364464: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M1
2025-04-03 08:55:48.364526: I metal_plugin/src/device/metal_device.cc:296]
systemMemory: 8.00 GB
2025-04-03 08:55:48.364575: I metal_plugin/src/device/metal_device.cc:313]
maxCacheSize: 2.67 GB
2025-04-03 08:55:48.364878: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2025-04-03 08:55:48.364900: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
```

<undefined>)

```
[19]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 96)	2,688
batch_normalization (BatchNormalization)	(None, 32, 32, 96)	384
max_pooling2d (MaxPooling2D)	(None, 15, 15, 96)	0
conv2d_1 (Conv2D)	(None, 15, 15, 256)	221,440
batch_normalization_1 (BatchNormalization)	(None, 15, 15, 256)	1,024
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 256)	0
conv2d_2 (Conv2D)	(None, 7, 7, 384)	885,120
batch_normalization_2 (BatchNormalization)	(None, 7, 7, 384)	1,536
conv2d_3 (Conv2D)	(None, 7, 7, 384)	1,327,488
batch_normalization_3 (BatchNormalization)	(None, 7, 7, 384)	1,536
conv2d_4 (Conv2D)	(None, 7, 7, 256)	884,992
batch_normalization_4 (BatchNormalization)	(None, 7, 7, 256)	1,024
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 256)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 512)	1,180,160
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328

dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2,570

Total params: 4,641,290 (17.71 MB)

Trainable params: 4,638,538 (17.69 MB)

Non-trainable params: 2,752 (10.75 KB)

```
[21]: history = model.fit(train_generator, epochs=10, validation_data=test_generator)
```

Epoch 1/10

/opt/anaconda3/envs/tf-metal/lib/python3.9/site-

packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:

UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self._warn_if_super_not_called()

2025-04-03 08:56:07.571112: I

tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:117]

Plugin optimizer for device_type GPU is enabled.

391/391 89s 221ms/step -

accuracy: 0.1971 - loss: 17.2317 - val_accuracy: 0.2390 - val_loss: 15.3230

Epoch 2/10

391/391 86s 220ms/step -

accuracy: 0.2729 - loss: 17.1955 - val_accuracy: 0.3836 - val_loss: 9.4347

Epoch 3/10

391/391 86s 220ms/step -

accuracy: 0.3409 - loss: 8.0069 - val_accuracy: 0.3791 - val_loss: 5.2993

Epoch 4/10

391/391 87s 223ms/step -

accuracy: 0.3934 - loss: 4.0883 - val_accuracy: 0.4176 - val_loss: 3.5327

Epoch 5/10

391/391 88s 225ms/step -

accuracy: 0.4287 - loss: 3.3027 - val_accuracy: 0.5277 - val_loss: 2.4025

Epoch 6/10

391/391 89s 228ms/step -

accuracy: 0.4528 - loss: 3.3160 - val_accuracy: 0.4090 - val_loss: 3.5003

Epoch 7/10

391/391 90s 230ms/step -

accuracy: 0.4615 - loss: 3.2472 - val_accuracy: 0.3516 - val_loss: 5.5601

Epoch 8/10

```

391/391          90s 230ms/step -
accuracy: 0.4846 - loss: 3.6987 - val_accuracy: 0.4589 - val_loss: 5.2175
Epoch 9/10
391/391          96s 247ms/step -
accuracy: 0.4819 - loss: 4.6660 - val_accuracy: 0.4294 - val_loss: 4.6699
Epoch 10/10
391/391          103s 263ms/step -
accuracy: 0.5187 - loss: 4.8957 - val_accuracy: 0.3013 - val_loss: 7.9245

```

```

[23]: y_test_labels=np.argmax(y_test, axis=1)
      y_pred = np.argmax(model.predict(x_test), axis=1)
      print("Classification Report:")
      print(classification_report(y_test_labels, y_pred))
      print("Confusion Matrix:")
      print(confusion_matrix(y_test_labels, y_pred))

```

```

313/313          5s 15ms/step
Classification Report:

```

	precision	recall	f1-score	support
0	0.21	0.97	0.34	1000
1	0.95	0.35	0.52	1000
2	0.15	0.51	0.23	1000
3	0.00	0.00	0.00	1000
4	0.45	0.01	0.03	1000
5	0.81	0.05	0.10	1000
6	0.97	0.16	0.27	1000
7	0.91	0.25	0.40	1000
8	0.62	0.26	0.37	1000
9	0.83	0.45	0.58	1000
accuracy			0.30	10000
macro avg	0.59	0.30	0.28	10000
weighted avg	0.59	0.30	0.28	10000

```

Confusion Matrix:
[[967  0 27  0  0  0  0  0  4  2]
 [458 355 12  0  0  1  0  1 90 83]
 [487  0 510  0  1  0  0  1  1  0]
 [309  0 666  0  2  8  2  5  4  4]
 [359  0 619  0 13  0  0  8  1  0]
 [259  0 677  0  0 51  3  6  2  2]
 [178  0 649  0  3  0 159  1 10  0]
 [488  0 243  0 10  2  0 253  1  3]
 [675  1  65  0  0  0  0  0 259  0]
 [457 18  33  0  0  1  0  2  43 446]]

```

```

/opt/anaconda3/envs/tf-metal/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:

```

Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/opt/anaconda3/envs/tf-metal/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/opt/anaconda3/envs/tf-metal/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[25]: precision = precision_score(y_test_labels, y_pred, average='weighted')
recall = recall_score(y_test_labels, y_pred, average='weighted')
f1 = f1_score(y_test_labels, y_pred, average='weighted')

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-score: {f1:.4f}")
```

Precision: 0.5894

Recall: 0.3013

F1-score: 0.2822

```
/opt/anaconda3/envs/tf-metal/lib/python3.9/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[27]: plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Model Accuracy')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Model Loss')
```

```
plt.show()
```

